

The Abundance Inversion — What Stays Expensive When Intelligence Gets Cheap

The Abundance Inversion

Operational Philosophy Brief — August 2026

In April, Elon Musk posted that AI and robotics will mean

"everyone can have a penthouse if they want,"

"sustainable abundance."



Strip the provocation and the argument is clean: if AI removes the barriers to entry in every sector, then every production cost collapses — software, video,

advice, design, marketing campaigns, drugs, research, education. We enter an economy of massive value creation sitting on top of rapid cost deflation.

Most reactions stop at the utopia or the dystopia. Both skip the mechanism. The interesting question is narrower and more useful: when a cost collapses, where does the value go? That question has a precise answer, and the answer tells you what to build, what to sell, and what to protect.

1 — Collapsing the cost of cognition

Every product and service is a stack of costs: **cognition** (design, drafting, analysis, judgment), **matter and energy**, **permission** (licenses, regulation, access), **proof** (testing, verification, audit), and **human hours**. Musk's claim becomes precise once you map it onto this stack: AI is repricing one layer — cognition — toward the cost of electricity. The other four layers are untouched by the model itself.

We have run this experiment before. The printing press collapsed the cost of copying text by orders of magnitude. The value of writing did not disappear; it migrated — from scribes to authors, from authors to publishers, from publishers to whoever controlled distribution and credibility. The internet repeated the pattern with distribution costs, and value migrated again, this time to whoever controlled attention.

Cost collapse never deletes value. It relocates value to the adjacent layer that is still scarce. Keep that rule; everything below is an application of it.

2 — Watching value migrate to the complements

The price curve is no longer speculative. Andreessen Horowitz has tracked a **"LLMflation"** trend: the cost of an equivalent-performance model falls by roughly 10x per year. The Stanford AI Index recorded GPT-3.5-level inference dropping from \$20 to \$0.07 per million tokens through 2024, and NVIDIA now advertises a Blackwell inference floor near **\$0.02 per million tokens**. Gartner forecasts a further **90% drop in inference costs by 2030**.

And yet enterprise AI bills keep rising. **Enterprise LLM spending tripled during 2025, and 93% of organizations report exceeding their AI budgets**, because agentic workflows consume 10 to 20 times more tokens per task. This is Jevons paradox applied to cognition: when a resource gets cheaper, total consumption grows faster than the price falls. Unit cost down, total spend up. The first correction to Musk's framing is that cheap intelligence and cheap total cost are different things.

The second mechanism matters more. Economists call it the complements rule: make one input abundant, and its complements capture the value. Free generation makes verification expensive. Free code makes operations, security, and integration expensive. Free content makes attention and provenance expensive. Abundance in one layer is a pricing event in every adjacent layer.

3 — Following the bottleneck across five sectors

The migration is already visible in the data. Each sector shows the same shape: the cognitive step deflates, and the bottleneck moves one layer down the stack.

Software. Roughly 41% of code is now AI-generated, and Anthropic's 2026 agentic coding data shows average agent sessions stretching from 4 to 23 minutes with dozens of autonomous tool calls. Building software is deflating toward the cost of compute. Meanwhile code churn has roughly doubled, and maintenance, security, and compliance dominate lifecycle cost. The moat moved from writing code to operating systems people trust.

Content and marketing. Supply of competent content is now effectively infinite, while human attention is fixed. The price of generic content converges to zero, and value concentrates in voice, provenance, and an audience that believes you. The flood of machine-generated material is what makes authenticity a paid feature.

Pharma. Isomorphic Labs received FDA clearance in January 2026 for ISM8969, among the first AI-designed drugs to enter human testing — with first trials now expected by end of 2026, a year behind the original promise. Discovery compressed; the trial did not. Patients, trial sites, and regulatory throughput are the bottleneck now. The scarce resource moved from the molecule to the proof.

Education. Content was already free before AI. The expensive layers were always motivation, feedback, and credentialing. AI tutors attack the feedback layer, but credentials survive because a degree sells trust in a person, and trust does not scale with tokens.

Advice. Legal drafts, tax memos, and consulting decks deflate toward zero. What persists is liability: a licensed professional who signs, and who loses something when wrong. The signature is the product.

The labor market is the cross-sector confirmation. Entry-level hiring in tech fell 30 to 50% in 2025, hiring dropped again in March 2026, and economists describe an "experience creep" where employers pay for judgment and

verification while machines do the generating. Even **Dario Amodei's stark entry-level warning** has been **tempered by payroll data**: aggregate employment holds, but the paid layer is shifting from producing work to vouching for it.

4 — Counting what stays scarce

Generalize the pattern and four categories survive the cost collapse. Together they form the scarcity stack of the abundance economy.

1. **Physics.** Energy, chips, land, grid connections. Intelligence has a floor price, and the floor is the electricity bill. Data center electricity demand **surged 17% in 2025** and is projected to roughly double to 945 TWh by 2030 — Japan's entire consumption. The five largest hyperscalers have committed **\$660–690 billion of capex for 2026**, and the binding constraint is **power delivery, with multi-year interconnection queues**.
2. **Proof.** Verification, provenance, security, audit. When generating an artifact costs nothing, knowing whether it is true, safe, and authentic becomes the expensive step.
3. **People.** Attention, time, care, taste, and legal accountability. Human hours do not scale, so they rise in relative price — the Baumol effect applied to the AI economy.
4. **Position.** Status, prime locations, incumbency, regulation, network effects — scarcity by construction, immune to productivity gains by definition.

The reversal

Here is the climax the cost-collapse story was hiding. Everyone asks: what will AI make cheap? The question that determines the next decade is the reverse: **what does cheap intelligence make expensive?**

The answer is the complements of intelligence — energy at the bottom of the stack, trust at the top. Abundance never abolishes scarcity; every abundance mints its own scarcity, and the mint is already running. The more content floods the zone, the more a verified word is worth. The more code agents ship, the more a trustworthy system costs. The more answers are free, the more a question worth asking becomes the asset.

Generation becomes worthless. Judgment becomes priceless. Musk's three questions — what does an abundance economy look like, how do you create

value, what stays rare — collapse into one: where is your bottleneck, and who owns it?

The price of being believed

Musk has the direction right and the shape incomplete. Costs collapse for anything made of cognition; value concentrates in anything made of proof. So run the audit this week. List what you sell. Cross out everything a model can generate. What remains — energy contracts, distribution, proprietary data, your name, your willingness to be accountable — that remainder is your business now.

For many of us, the remainder is smaller than we thought. The last scarce resource in the abundance economy is not intelligence. It is being believed. And an economy that floods the world with free answers is quietly repricing every honest one.

The control set: running your own scarcity audit

Philosophy that operates. Four questions for your next staff meeting:

1. **List** what you sell, at the level of individual deliverables.
2. **Cross out** every deliverable a model can generate at acceptable quality.
3. **Name** the scarcity behind what remains: physics, proof, people, or position.
4. **Decide** one investment that strengthens a remaining layer — and stop funding a layer that is deflating.

If nothing remains after step two, you have a product, not a business. The same applies to a career.

Raphaël MANSUY